

TF-IDF Cosine Baseline vs. LLM-as-a-Judge
(same 40-case test set, same bucket thresholds: ≥ 0.65 correct, ≥ 0.30 partial)

